

# CS 588 Assignment 5, Part 1

Bhavya Chawla (A20516957) [bchawla@hawk.iit.edu](mailto:bchawla@hawk.iit.edu)

## (PART I: 150 points) **Requirements Specification:**

1. Use **Go** languages, **Docker**, **GCP** (Google Cloud Platform), **Postgres** to create microservices to collect data of the **Posts** (Question/body and Answer/body) from **StackOverflow** and **Issues** from **GitHub** and store the data collected in two databases: **StackoverflowDB** and **GitHubDB** using **Postgres** database server.
2. Collect the data from **StackOverflow** and **GitHub** for the following frameworks/libraries:
  - a. Prometheus
    - i. <https://stackoverflow.com/search?q=Prometheus>
    - ii. <https://github.com/prometheus/prometheus>
  - b. Selenium
  - c. OpenAI
  - d. Docker
  - e. Milvus
  - f. Go
    - i. <https://stackoverflow.com/search?q=golang>
    - ii. <https://github.com/golang/go>
3. You will need to create **developer accounts** and **API KEYS** on **GitHub** and **StackExchange**
  - <https://api.github.com>
  - <https://docs.github.com/en>
  - Here is the URL to guide you on how to generate your GITHUB\_TOKEN
    - <https://help.github.com/articles/creating-an-access-token-for-command-line-use/>
  - Create your GitHub personal access tokens from
    - <https://github.com/settings/tokens>
  - <https://api.stackexchange.com/>
  - <https://api.stackexchange.com/docs>
  - <https://stackapps.com/apps/oauth/register>
4. Use Prometheus to show how many API calls to GitHub and StackOverflow were made per second, and how much data collected per second.
5. Run 3 experiments using data collected for past 2 days, 7 days, and 45 days.

# Overview:

This project aims to collect data from StackOverflow and GitHub related to specific frameworks/libraries and store the information in two separate databases, StackoverflowDB and GitHubDB, using Postgres as the database server. The chosen frameworks/libraries are Prometheus, Docker, and Go.

## Technologies Used:

- Programming Language: Go
- Containerization: Docker
- Cloud Platform: Google Cloud Platform (GCP)
- Database: Postgres
- Monitoring Tool: Prometheus

## Steps:

### 1. Created Developer Accounts and API Keys

GitHub

- Generated a GitHub personal access token
- Created a personal access token on GitHub.

StackOverflow

- Obtained the necessary API key.

### 2. Running the Microservices

Step 1: Clone the Repository

bash

Copy code

```
git clone https://github.com/your-username/stackoverflow-github-microservices.git  
cd src
```

Step 2: Build and Run Docker Containers

```
docker-compose up --build
```

### 3. Data Collection

The microservices will start collecting data from StackOverflow and GitHub for the specified frameworks/libraries. The data will be stored in the StackoverflowDB and GitHubDB databases on the Postgres server.

## 4. Prometheus Monitoring

Prometheus is configured to monitor API calls to GitHub and StackOverflow. Metrics are shown at the end of the report

## 5. Run Experiments

Three experiments were conducted using data collected over the past 2 days, 7 days, and 45 days.

### Metrics:

```
# HELP github_api_calls_2_days_total Total number of API calls made to GitHub in the past 2 days
# TYPE github_api_calls_2_days_total counter
github_api_calls_2_days_total 762
# HELP github_api_calls_45_days_total Total number of API calls made to GitHub in the past 45 days
# TYPE github_api_calls_45_days_total counter
github_api_calls_45_days_total 546
# HELP github_api_calls_7_days_total Total number of API calls made to GitHub in the past 7 days
# TYPE github_api_calls_7_days_total counter
github_api_calls_7_days_total 623
# HELP github_api_calls_per_second Rate of API calls made to GitHub per second
# TYPE github_api_calls_per_second summary
github_api_calls_per_second_sum{endpoint="github_endpoint"} 7.856890268377948
github_api_calls_per_second_count{endpoint="github_endpoint"} 18
# HELP go_gc_duration_seconds A summary of the pause duration of garbage collection cycles.
# TYPE go_gc_duration_seconds summary
go_gc_duration_seconds{quantile="0"} 6.3042e-05
go_gc_duration_seconds{quantile="0.25"} 0.0001115
go_gc_duration_seconds{quantile="0.5"} 0.000134416
go_gc_duration_seconds{quantile="0.75"} 0.000229501
go_gc_duration_seconds{quantile="1"} 0.000559417
go_gc_duration_seconds_sum 0.008311496
go_gc_duration_seconds_count 44
```

```

# HELP promhttp_metric_handler_requests_in_flight current number of scrapes being served.
# TYPE promhttp_metric_handler_requests_in_flight gauge
promhttp_metric_handler_requests_in_flight 1
# HELP promhttp_metric_handler_requests_total Total number of scrapes by HTTP status code.
# TYPE promhttp_metric_handler_requests_total counter
promhttp_metric_handler_requests_total{code="200"} 9
promhttp_metric_handler_requests_total{code="500"} 0
promhttp_metric_handler_requests_total{code="503"} 0
# HELP stackoverflow_api_calls_2_days_total Total number of API calls made to StackOverflow in the past 2 days
# TYPE stackoverflow_api_calls_2_days_total counter
stackoverflow_api_calls_2_days_total 0
# HELP stackoverflow_api_calls_45_days_total Total number of API calls made to StackOverflow in the past 45 days
# TYPE stackoverflow_api_calls_45_days_total counter
stackoverflow_api_calls_45_days_total 0
# HELP stackoverflow_api_calls_7_days_total Total number of API calls made to StackOverflow in the past 7 days
# TYPE stackoverflow_api_calls_7_days_total counter
stackoverflow_api_calls_7_days_total 0
# HELP stackoverflow_api_calls_per_second Rate of API calls made to StackOverflow per second
# TYPE stackoverflow_api_calls_per_second summary
stackoverflow_api_calls_per_second_sum{endpoint="stackoverflow_endpoint"} 343.2951616792917
stackoverflow_api_calls_per_second_count{endpoint="stackoverflow_endpoint"} 18
# HELP total_github_api_calls Total number of API calls made to GitHub
# TYPE total_github_api_calls counter
total_github_api_calls 18
# HELP total_stackoverflow_api_calls Total number of API calls made to StackOverflow
# TYPE total_stackoverflow_api_calls counter
total_stackoverflow_api_calls 18

```

```

# HELP github_api_calls_2_days_total Total number of API calls made to GitHub in the past 2 days
# TYPE github_api_calls_2_days_total counter
github_api_calls_2_days_total 762
# HELP github_api_calls_45_days_total Total number of API calls made to GitHub in the past 45 days
# TYPE github_api_calls_45_days_total counter
github_api_calls_45_days_total 546
# HELP github_api_calls_7_days_total Total number of API calls made to GitHub in the past 7 days
# TYPE github_api_calls_7_days_total counter
github_api_calls_7_days_total 623
# HELP github_api_calls_per_second Rate of API calls made to GitHub per second
# TYPE github_api_calls_per_second summary
github_api_calls_per_second_sum{endpoint="github_endpoint"} 7.856890268377948
github_api_calls_per_second_count{endpoint="github_endpoint"} 18
# HELP go_gc_duration_seconds A summary of the pause duration of garbage collection cycles.
# TYPE go_gc_duration_seconds summary
go_gc_duration_seconds{quantile="0"} 6.3042e-05
go_gc_duration_seconds{quantile="0.25"} 0.0001115
go_gc_duration_seconds{quantile="0.5"} 0.000134416
go_gc_duration_seconds{quantile="0.75"} 0.000229501
go_gc_duration_seconds{quantile="1"} 0.000559417
go_gc_duration_seconds_sum 0.008311496
go_gc_duration_seconds_count 44
# HELP go_goroutines Number of goroutines that currently exist.
# TYPE go_goroutines gauge
go_goroutines 12
# HELP go_info Information about the Go environment.

```

```
# TYPE go_info gauge
go_info{version="go1.21.4"} 1
# HELP go_memstats_alloc_bytes Number of bytes allocated and still in use.
# TYPE go_memstats_alloc_bytes gauge
go_memstats_alloc_bytes 1.9816e+06
# HELP go_memstats_alloc_bytes_total Total number of bytes allocated, even if freed.
# TYPE go_memstats_alloc_bytes_total counter
go_memstats_alloc_bytes_total 7.7657376e+07
# HELP go_memstats_buck_hash_sys_bytes Number of bytes used by the profiling bucket hash table.
# TYPE go_memstats_buck_hash_sys_bytes gauge
go_memstats_buck_hash_sys_bytes 5191
# HELP go_memstats_frees_total Total number of frees.
# TYPE go_memstats_frees_total counter
go_memstats_frees_total 301514
# HELP go_memstats_gc_sys_bytes Number of bytes used for garbage collection system metadata.
# TYPE go_memstats_gc_sys_bytes gauge
go_memstats_gc_sys_bytes 4.47256e+06
# HELP go_memstats_heap_alloc_bytes Number of heap bytes allocated and still in use.
# TYPE go_memstats_heap_alloc_bytes gauge
go_memstats_heap_alloc_bytes 1.9816e+06
# HELP go_memstats_heap_idle_bytes Number of heap bytes waiting to be used.
# TYPE go_memstats_heap_idle_bytes gauge
go_memstats_heap_idle_bytes 3.743744e+06
# HELP go_memstats_heap_inuse_bytes Number of heap bytes that are in use.
# TYPE go_memstats_heap_inuse_bytes gauge
go_memstats_heap_inuse_bytes 4.05504e+06
# HELP go_memstats_heap_objects Number of allocated objects.
# TYPE go_memstats_heap_objects gauge
go_memstats_heap_objects 5545
# HELP go_memstats_heap_released_bytes Number of heap bytes released to OS.
# TYPE go_memstats_heap_released_bytes gauge
go_memstats_heap_released_bytes 1.490944e+06
# HELP go_memstats_heap_sys_bytes Number of heap bytes obtained from system.
# TYPE go_memstats_heap_sys_bytes gauge
go_memstats_heap_sys_bytes 7.798784e+06
# HELP go_memstats_last_gc_time_seconds Number of seconds since 1970 of last garbage collection.
# TYPE go_memstats_last_gc_time_seconds gauge
go_memstats_last_gc_time_seconds 1.7012269834825652e+09
# HELP go_memstats_lookups_total Total number of pointer lookups.
# TYPE go_memstats_lookups_total counter
go_memstats_lookups_total 0
# HELP go_memstats_mallocs_total Total number of mallocs.
# TYPE go_memstats_mallocs_total counter
go_memstats_mallocs_total 307059
```

```
# HELP go_memstats_mcache_inuse_bytes Number of bytes in use by mcache structures.
# TYPE go_memstats_mcache_inuse_bytes gauge
go_memstats_mcache_inuse_bytes 9600
# HELP go_memstats_mcache_sys_bytes Number of bytes used for mcache structures obtained from
system.
# TYPE go_memstats_mcache_sys_bytes gauge
go_memstats_mcache_sys_bytes 15600
# HELP go_memstats_mspan_inuse_bytes Number of bytes in use by mspan structures.
# TYPE go_memstats_mspan_inuse_bytes gauge
go_memstats_mspan_inuse_bytes 161952
# HELP go_memstats_mspan_sys_bytes Number of bytes used for mspan structures obtained from
system.
# TYPE go_memstats_mspan_sys_bytes gauge
go_memstats_mspan_sys_bytes 179256
# HELP go_memstats_next_gc_bytes Number of heap bytes when next garbage collection will take place.
# TYPE go_memstats_next_gc_bytes gauge
go_memstats_next_gc_bytes 4.194304e+06
# HELP go_memstats_other_sys_bytes Number of bytes used for other system allocations.
# TYPE go_memstats_other_sys_bytes gauge
go_memstats_other_sys_bytes 1.654705e+06
# HELP go_memstats_stack_inuse_bytes Number of bytes in use by the stack allocator.
# TYPE go_memstats_stack_inuse_bytes gauge
go_memstats_stack_inuse_bytes 589824
# HELP go_memstats_stack_sys_bytes Number of bytes obtained from system for stack allocator.
# TYPE go_memstats_stack_sys_bytes gauge
go_memstats_stack_sys_bytes 589824
# HELP go_memstats_sys_bytes Number of bytes obtained from system.
# TYPE go_memstats_sys_bytes gauge
go_memstats_sys_bytes 1.471592e+07
# HELP go_threads Number of OS threads created.
# TYPE go_threads gauge
go_threads 12
# HELP promhttp_metric_handler_requests_in_flight Current number of scrapes being served.
# TYPE promhttp_metric_handler_requests_in_flight gauge
promhttp_metric_handler_requests_in_flight 1
# HELP promhttp_metric_handler_requests_total Total number of scrapes by HTTP status code.
# TYPE promhttp_metric_handler_requests_total counter
promhttp_metric_handler_requests_total{code="200"} 9
promhttp_metric_handler_requests_total{code="500"} 0
promhttp_metric_handler_requests_total{code="503"} 0
# HELP stackoverflow_api_calls_2_days_total Total number of API calls made to StackOverflow in the
past 2 days
# TYPE stackoverflow_api_calls_2_days_total counter
stackoverflow_api_calls_2_days_total 0
```

```
# HELP stackoverflow_api_calls_45_days_total Total number of API calls made to StackOverflow in the
past 45 days
# TYPE stackoverflow_api_calls_45_days_total counter
stackoverflow_api_calls_45_days_total 0
# HELP stackoverflow_api_calls_7_days_total Total number of API calls made to StackOverflow in the
past 7 days
# TYPE stackoverflow_api_calls_7_days_total counter
stackoverflow_api_calls_7_days_total 0
# HELP stackoverflow_api_calls_per_second Rate of API calls made to StackOverflow per second
# TYPE stackoverflow_api_calls_per_second summary
stackoverflow_api_calls_per_second_sum{endpoint="stackoverflow_endpoint"} 343.2951616792917
stackoverflow_api_calls_per_second_count{endpoint="stackoverflow_endpoint"} 18
# HELP total_github_api_calls Total number of API calls made to GitHub
# TYPE total_github_api_calls counter
total_github_api_calls 18
# HELP total_stackoverflow_api_calls Total number of API calls made to StackOverflow
# TYPE total_stackoverflow_api_calls counter
total_stackoverflow_api_calls 18
```